Cody Lee

cody.lee.cl1@gmail.com • (626) 838-4155 • GitHub • Portfolio • Open to relocation

AI agent systems engineer who builds evaluation-driven, tool-using systems and measures where they break — across **measurement** (evaluation frameworks), **execution** (autonomous agents), and **adversarial** testing.

Education

University of California, Riverside • B.S. Computer Science — Jun 2026

Relevant coursework: Natural Language Processing, Artificial Intelligence, Data Science, Probability & Statistics

Experience

Cofounding Member — Rebase/Yield DeFi Protocol — 2021–2022

- Designed the protocol's token mechanism (rebase/yield model) and core economic incentives; the system attracted ∼$100M in total value locked at peak.
- Oversaw protocol development end to end and led investor onboarding, including angel-round relationships and early-contributor ramp.

Projects

Meridian: Forensic RAG Framework + Vector DB Benchmark 8 page research — May–Jul 2026

- Built a framework that measures whether an AI document-search system actually works — rigorous enough to catch its own measurement bugs. Tested on ∼800 questions over ∼79M characters of legal contracts.
- Two independent scoring layers (deterministic, plus a separate AI-judged one) keep a real gain from being mistaken for a measurement glitch; the tuned system lifted contract answer accuracy from 26% to 75% and transferred to medical data with no retuning.
- Extended the same rigor to a live turbopuffer vs Qdrant vector-DB benchmark over 2M passages on a co-located GCP VM — measuring latency (warm/cold), recall@10, and cost: turbopuffer ∼15× cheaper at multi-tenant scale, Qdrant ∼2× faster warm, with a documented audit of the measurement bugs caught.

Aether: GTM Agent Platform + LLM Reasoning Engine Live demo — Apr–Jun 2026

- Triages inbound leads and runs outbound account campaigns on one reason/act/observe loop — parsing, enriching, scoring, and drafting per lead, and expanding accounts into ICP-scored lookalike targets.
- Validated on FinQA: 75.5% end to end with retrieval 85% source-accurate; cites its source or refuses rather than guessing. Direct SDK, no LangChain.
- Grounds every fact across People Data Labs, Apollo, and Brave and syncs live demand with Productboard. Built on FastAPI, Next.js, and Postgres (Neon) with Langfuse tracing; deployed on Vercel and Render.

Aegis: Deterministic Verifier + Eval Harness Verifier | Infra | Agent — Jun 2026

- Deterministic verifier that judges whether a patch genuinely closes a bug instead of just passing tests; validated on real CVEs (MLflow, LibreChat) on live Docker targets — caught a maintainer's fix that passes the full suite yet leaves the bug open.
- Built a benchmark-agnostic parallel eval harness (CVE-Bench on Inspect) and ran a controlled bare-vs-scaffold study that isolated the real bottleneck instead of chasing a number.

Polymarket Trading Bot Autopsy 2 page summary | 15 page autopsy — Apr to Jun 2026

- Built, ran, then dissected a trading-bot system: found three hidden measurement bugs that had inflated backtests ∼135×, so the best-looking bots lost the most real money.
- Classified thousands of trader wallets with a tiered model pipeline (cheap first, escalate only when needed) — ∼4× cheaper across 180 bots and 417K simulated trades.

Technical Skills

Languages Python, TypeScript, SQL, Bash, Next.js
LLM / RAG Qdrant, Pinecone, LangGraph, MCP servers, Anthropic / OpenAI / Google SDKs, Pydantic v2
GTM / FDE HubSpot CRM, Apollo, Productboard, Brave Search, People Data Labs, n8n, Vercel / Render
Data & Observability pandas, DuckDB, SQLite, PostgreSQL (Neon); Langfuse, OpenTelemetry, Sentry
Infrastructure Git, Docker, Google Cloud (Compute Engine), Linux / WSL2, REST APIs, FastAPI, Streamlit