

Cody Lee

cody.lee.cl1@gmail.com • (626) 838-4155 • [GitHub](#) • [Portfolio](#) • Open to relocation

AI agent systems engineer building evaluation-driven, tool-using systems and analyzing where they fail. Work spans measurement, execution, and adversarial stress testing.

Education

University of California, Riverside • B.S. Computer Science

Expected Jun 2026

Relevant coursework: Natural Language Processing, Artificial Intelligence, Data Science, Probability & Statistics

Projects

Meridian — Forensic RAG Measurement Framework [8-page research](#)

May 2026

- Built a **forensic evaluation framework** for retrieval systems — rigorous enough to **catch its own measurement bugs**. Validated on a large legal benchmark: **~800 questions** across **~79M characters** of contracts, NDAs, and privacy policies.
- Split scoring into two independent layers — a deterministic one with **no AI grading AI**, plus a separate AI-judged layer for answer quality — so a genuine improvement is never mistaken for a measurement glitch.
- The tuned system retrieved **2–20× more accurately** than the published industry baseline and lifted answer accuracy on contracts from **26% to 75%**; the approach also transferred to a medical dataset with no retuning.

Aether — LLM Workflow Engine

Apr 2026

- Built a **reason-act-observe agent** that queries financial documents via SQL and produces **grounded charts or answers where every value is traceable to a source cell** — returns an evidence-backed refusal when the document can't support the question.
- **75.5% end-to-end on FinQA** across **200 questions**; retrieval hit the right source in **85%** of cases.
- Every reasoning step, SQL query, and tool call is logged in a full trace — a reviewer can see exactly which row produced which number in the chart. Direct SDK, no LangChain.

Polymarket Trading Bot Autopsy [2-page summary](#) | [15-page autopsy](#)

Apr–Jun 2026

- Identified **three measurement bugs** in a trading bot that inflated backtests **~135×** — reversing real-world performance rankings.
- Classified thousands of trader wallets with a tiered AI pipeline that tried cheap models first and escalated only when needed — **~4× cheaper** while running 180 bots across **417K simulated trades**.

ChainTax — Crypto Tax Engine

Feb–Mar 2026

- Built a cross-source crypto reconciliation pipeline — **540K+** events across exchanges and blockchains — that constructs a unified ledger and prevents cross-chain double-counting.
- Generates audit-ready IRS tax forms under **three competing accounting methods**, so a user can compare outcomes side by side.

Aegis — Verifiable RL Environment for Defensive Cybersecurity

In progress

- Building a training environment for AI agents that find and fix software vulnerabilities, with a **tamper-resistant verifier** designed to detect **reward hacking** — agents optimizing the checker instead of the bug.
- Evaluating it two ways: how capable the agent is on standard security benchmarks, and how reliably the checker catches gamed fixes — applying **Meridian's measurement rigor** to AI training rewards.

Technical Skills

Languages Python, TypeScript, SQL, Bash

LLM / RAG Hybrid retrieval, chunking, embedding, reranking, fusion, deterministic span-overlap evaluation, vector databases (Qdrant, Pinecone), LangGraph, MCP servers, Anthropic / OpenAI / Google SDKs

Data & Observability Cross-source ingestion, pandas, DuckDB, SQLite; Langfuse tracing, cost / latency

Infrastructure Git, REST APIs, FastAPI, Streamlit, VPS deployment

Methodology Forensic measurement, failure-mode analysis, statistical evaluation, technical documentation