

Cody Lee

cody.lee.cl1@gmail.com • (626) 838-4155 • [GitHub](#) • [Portfolio](#) • Open to relocation

AI agent systems engineer building evaluation driven, tool using systems and measuring where they break. Work spans three layers: **measurement**, deterministic evaluation frameworks; **execution**, autonomous document agents; and **adversarial**, stress testing systems and their reward signals.

Education

University of California, Riverside • B.S. Computer Science

Expected Jun 2026

Relevant coursework: Natural Language Processing, Artificial Intelligence, Data Science, Probability & Statistics

Projects

Meridian: Forensic RAG Measurement Framework [8 page research](#)

May 2026

- Built a framework that measures whether an AI document search system actually works, rigorous enough to catch its own measurement bugs. Tested on a large legal benchmark: ~800 questions across ~79M characters of contracts, NDAs, and privacy policies.
- Split scoring into two independent layers: a deterministic one with no AI grading AI, plus a separate AI judged layer for answer quality, so a genuine improvement is never mistaken for a measurement glitch.
- The tuned system retrieved 2 to 20× more accurately than the published industry baseline and lifted answer accuracy on contracts from 26% to 75%; the approach also transferred to a medical dataset with no retuning.

Aether: LLM Workflow Engine

Apr 2026

- Built a reason/act/observe agent for financial documents: give it a fund statement, P&L CSV, or compliance PDF and a plain English goal, and it queries the data with SQL and returns a grounded chart or answer where every number traces back to a source cell, or refuses, with evidence, when the document can't support it.
- 75.5% end to end on FinQA across 200 questions; retrieval hit the right source in 85% of cases.
- Every reasoning step, SQL query, and tool call is logged in a full trace, so a reviewer can see exactly which row produced which number in the chart. Direct SDK, no LangChain.

Polymarket Trading Bot Autopsy [2 page summary](#) | [15 page autopsy](#)

Apr to Jun 2026

- Built, ran, and then dissected a systematic trading bot project: three hidden measurement bugs had inflated backtest results ~135×, so the bots that looked best on paper performed worst with real money.
- Classified thousands of trader wallets with a tiered AI pipeline that tried cheap models first and escalated only when needed, ~4× cheaper while running 180 bots across 417K simulated trades.

ChainTax: Crypto Tax Engine

Feb to Mar 2026

- Built a pipeline that pulls 540K+ crypto transactions from multiple exchanges and blockchains, reconciles them into one clean ledger, and detects transfers to avoid double counting.
- Generates audit ready IRS tax forms under three competing accounting methods, so a user can compare outcomes side by side.

Aegis: Verifiable RL Environment for Defensive Cybersecurity

In progress

- Building a training environment where an AI agent learns to find and fix real software vulnerabilities, graded by a tamper resistant automated checker built to catch agents that game the reward, patching the test instead of the actual bug.
- Evaluating it two ways: how capable the agent is on standard security benchmarks, and how reliably the checker catches gamed fixes, applying Meridian's measurement rigor to AI training rewards.

Technical Skills

Languages Python, TypeScript, SQL, Bash

LLM / RAG Hybrid retrieval, chunking, embedding, reranking, fusion, deterministic span overlap evaluation, vector databases (Qdrant, Pinecone), LangGraph, MCP servers, Anthropic / OpenAI / Google SDKs

Data & Observability Cross source ingestion, pandas, DuckDB, SQLite; Langfuse tracing, cost / latency

Infrastructure Git, REST APIs, FastAPI, Streamlit, VPS deployment

Methodology Forensic measurement, failure mode analysis, statistical evaluation, technical documentation