

CODY LEE

Los Angeles, CA · Open to relocation
cody.lee.cl1@gmail.com · (626) 838-4155
GitHub: github.com/clee12111 · Portfolio: codylee.tech

Building end-to-end LLM pipelines — retrieval, agent loops, evaluation — and forensically documenting where they break.

EDUCATION

University of California, Riverside · B.S. Computer Science

Expected Jun 2026

Relevant coursework: Natural Language Processing, Artificial Intelligence, Data Science Ethics, Probability & Statistics, Operating Systems

PROJECTS

Production RAG Stack Forensics

2026

Published: 5-page autopsy → github.com/clee12111/RagForensics

- Built a **production RAG system** on the industry-standard stack — **LangGraph**, **Pinecone**, **Langfuse**, and **cross-provider inference** (Claude, GPT-5.5, Gemini) — and documented **four failure modes** with measured fixes against a **150-question eval suite**.
- Caught an **evaluation bug** scoring answers against empty retrieved content — hidden until a different model exposed it — that had **inverted a result**; the fix reversed it, preventing a wrong published finding.
- Showed model capability wasn't the bottleneck with strong retrieval: **0.23 quality spread across a 28× cost difference** between providers.
- Built a custom **MCP server** exposing eval, latency, cost, and failure-mode data as **agent-queryable tools** over the open Model Context Protocol.

Polymarket Trading Bot Autopsy

Apr–Jun 2026

Published: 2-page summary → github.com/clee12111/polymarket-autopsy

- Found **three measurement bugs** that inflated paper trading results **~135×** and made the strategy misleading — the bots flagged as best performed worst live.
- Built a **3-layer LLM classification pipeline** (Haiku → Sonnet → Opus) over **2,409 wallets**; 180 bots, 417K simulated trades. **Tiered routing cut cost** to \$30–40 vs ~\$150 on Opus alone.
- Operated a **live execution environment** with wallet auth, kill switches, and trade logging; reconciled millions of cross-source transactions across **Polymarket, Binance, and on-chain data**.

Aether — LLM Workflow Engine

Apr 2026

- Built a **hybrid RAG pipeline** over compliance documents — BM25 + dense embeddings fused via **Reciprocal Rank Fusion**, reranked with a **cross-encoder**; **planner / executor / critic agent loop** with audit trails.
- 96% retrieval precision** (reproducible across 5 runs), **87% end-to-end correctness**, **\$0.65 per run** via tiered model routing.

ChainTax — Crypto Tax Engine

Feb–Mar 2026

- Built a **cross-source ingestion and reconciliation pipeline**: **540K+ events** across Hyperliquid and two on-chain APIs (Alchemy, DeBank), with bridge detection to prevent double-counting across chains.
- FIFO lot matching** with three competing IRS funding treatments producing three **Form 8949 variants** per run, plus offsetting-position detection across spot, perp, and correlated pairs.

TECHNICAL SKILLS

Languages	Python, TypeScript, SQL, Bash
LLM / RAG	Hybrid retrieval (dense + sparse), chunking, embedding, reranking, vector databases (Pinecone, ChromaDB), prompt engineering , LLM evaluation , agent orchestration (LangGraph), MCP servers , Anthropic / OpenAI / Google SDKs
Observability	Langfuse tracing, per-stage cost / latency instrumentation
Data	Cross-source ingestion and reconciliation, pandas, DuckDB, SQLite, async websocket pipelines
Infrastructure	Git, REST APIs, FastAPI, Streamlit, VPS deployment
Methodology	Forensic measurement, failure-mode analysis, statistical evaluation, technical documentation